

Homework - 7

Maximum points = 20 (10+10)

Results to be uploaded in Blackboard

Question-1

Match each of these rows to MongoDB or Cassandra or Both or N/A (not applicable)

1. Distributed NoSQL Database
2. All nodes in clusters are peers
3. Highly Available DB
4. Highly Consistent DB
5. Tunable Consistency

Sample answer will look like :

Distributed NoSQL Database = N/A

This sample answer may or may not be true

Question-2

Fill in the blanks

Query-First Design is applicable to ___ DB. This means that instead of looking at the data or domain model and designing the ___ as done in RDBMS, we design the ___ first based on application demands. Then based on the data query needs, define the ___ structure.

Queries on OLTP databases typically process ___ amount of rows whereas analytical queries on data warehouses process ___ amounts of rows.

In a Key-Value store DB, we query based on the ___ and this returns the ___ .

Data lake is more than just a dump of data in ___ storage. If there is lack of proper organization, data lakes degenerate into data ___ .